

Constraint Satisfaction Problems (CSP)

- Reading: Russell and Norvig, ch. 6
 - The material here is an edited version of Russell's slides
- Intro: map coloring example
- Domains
- Varieties of constraints
- Algorithms
 - Backtracking search

CSP intro

- Can use CSP when:
 - State defined as set of variables with specific values
 - Check for goal state by checking if given combination of values is acceptable
 - By limiting problem, supports different algorithms
- For N-Queens,
 - What would the state be?
 - How would you check if a state is a goal state?
- 8-puzzle doesn't really work because the steps are part of the solution
 - Cannot just swap positions of tiles

Example: Map coloring

- Goal: color areas (ex: provinces, countries) of a map so that adjacent areas always have different colors
- State: color for each area
- Can convert to a graph:
 - Each area is a vertex
 - Connect vertex to vertices representing adjacent areas
 - Example of *binary* CSP: constraint relates at most 2 variables
- Graph theory result: need at most 4 colors
 - see, for example: <https://www.mathsisfun.com/activity/coloring.html>

Example: Map coloring

- https://en.wikipedia.org/wiki/Graph_coloring#Applications :
 - Scheduling:
 - vertices represent jobs
 - edges represent jobs that conflict – ex: require same resource or same person
 - “coloring” = assigning time slot to job
 - Register allocation:
 - When compiler generates code, more efficient to store program variables in registers
 - If program variables needed at same time, cannot use same register
 - “coloring” = assigning register to program variable
 - C/C++ `register` keyword is a hint to use a register, if possible

Domain of problem variables

- Discrete
 - In general, harder than continuous domain!
 - Many problems for even finite domains are NP-Complete
 - Ex: Boolean satisfiability
 - In infinite domains (integers, strings)
 - Linear constraints are still solvable
 - Nonlinear constraints lead to undecidable problems
- Continuous
 - Linear constraints: solvable in polynomial time by linear programming algorithms
 - Fortunately, real world problems tend to be continuous

Varieties of constraints

- Unary: constraints on only one variable
 - Ex: vertex for California must be colored gold
- Binary: constraints apply to at most 2 variables
 - Ex: map coloring
- Higher order:
 - Ex: Cryptarithmic column constraints (next example)
- *Preferences*: soft constraints
 - Ex: try to schedule different sections of Calc 1 on MWF and TuTh
 - Often modeled as costs for assignments that violate constraints

Cryptarithmic example

- Constraints defined by equation with numbers encoded by words
- Each letter represents a distinct digit (0-9)
 - Goal: assign digits to letters to get valid equation
- Ex:

$$\begin{array}{r} \text{TWO} \\ + \text{TWO} \\ \hline \text{FOUR} \end{array}$$

Algorithms

- Start with straightforward, dumb approach
- Depth-first search
 - Assign value to unassigned variable that does not conflict with existing assignments
 - Backtrack if cannot find valid assignment for variable
 - Fail if failure with assigning values in all permutations of variables
- Branching factor = $(n-L)d$ for n variables at level L , domain size = d
 - $n!d^n$ leaves: can assign variables in any order, exponential in size of domain

Backtracking search

- Variable assignments are *commutative* if swapping values still satisfies constraints
 - Ex: for map coloring, solution groups areas into sets with same color
 - Can use any permutation of colors
- Branching factor reduced to d
 - At each level, only consider one variable

Backtracking search pseudocode

function Backtracking-Search(csp) returns solution/failure

 return Recursive-Backtracking({},csp)

function Recursive-Backtracking(assignment,csp) returns soln/failure

 if assignment is complete then return assignment

 var←Select-Unassigned-Variable(Variables[csp],assignment,csp)

 for each value in Order-Domain-Values(var,assignment,csp) do

 if value is consistent with assignment given Constraints[csp] then

 add {var = value} to assignment

 result←Recursive-Backtracking(assignment,csp)


 if result ≠ failure then return result

 remove {var = value} from assignment

 return failure

Backtracking example

- Edited version of Russell's figure:
 - (J. Neal Richter)
 - Bottom two leaf nodes swapped
 - What happens at the bottom left?



Backtracking efficiency

- General methods can greatly improve efficiency
- To consider:
 - Which variable should be assigned next?
 - In what order should its values be tried?
 - Can we detect inevitable failure early?
 - Can we take advantage of problem structure?

Minimum remaining values

- In backtracking example, assign to South Australia (bottom center) before Queensland (top right) because 2 of its neighbors have already been given a color
- Sudoku row
 - Obviously, if 8 columns assigned, last column only has 1 legal value left
 - If 7 columns assigned, assigning to last 2 columns may require backtracking, but just once

Degree heuristic

- Tiebreaker for minimum remaining value variables
 - Choose variable with the most constraints on remaining variables
 - Ex: for map coloring, the area with the most neighbors that remain uncolored

Least constraining value


- Choose value that rules out the fewest values in remaining variables
- In Cryptarithmic example ($TWO + TWO = FOUR$):
 - We know $F = 1$
 - O, U, R must be even
 - Therefore, try to pick odd values for T and W
- Using minimum remaining value, the degree heuristic and least constraining value, can solve N-queens problem for problem with N about 40 times greater

Early detection of failure

- Forward checking
 - Keep track of remaining legal values for each variable
 - Can backtrack as soon as any variable has no legal values
 - Ex: In Sudoku, if guesses lead to row only missing a 1, but with a 1 in same column as empty square

Early detection of failure

- Forward checking does not detect all inevitable failures
- Ex: Northern Territory and South Australia both can only be blue



Arc consistency

- For arc (edge) from X to Y ,
 - Arc is consistent iff for every value of X , there is at least one legal value of Y
- If variable assignment removes any value from X , X 's neighbors need to be rechecked
 - Ex: for coloring Australia, making Queensland green removes that color option from New South Wales

Arc consistency algorithm

function AC-3(csp) returns the CSP, possibly with reduced domains

inputs: csp, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$

local variables: queue, a queue of arcs, initially all the arcs in csp

while queue is not empty do

$(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$

 if Remove-Inconsistent-Values(X_i, X_j) then

 for each X_k in Neighbors[X_i] do

 add (X_k, X_i) to queue

function Remove-Inconsistent-Values(X_i, X_j) returns true iff succeeds

 removed \leftarrow false

 for each x in Domain[X_i] do

 if no value y in Domain[X_j] allows (x, y) to satisfy the constraint $X_i \leftrightarrow X_j$ then

 delete x from Domain[X_i]; removed \leftarrow true

return removed

Problem structure

- For coloring Australia,
 - Tasmania is an island, so can consider it separately from other territories
 - In general, disconnected parts of constraint graph allow them to be considered separately
- In general, if possible to divide problem into subproblems, each with c of the variables, can reach linear cost in n ($n/c * d^c$)
 - For $n=80$, $d=2$, $c=20$,
 - Execution time goes from 4 billion years to .4 seconds

Tree-structured problems

- If no cycles in constraints, problem is solvable in $O(n d^2)$ vs $O(d^n)$ for general CSP
- Algorithm:
 1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering
 2. For j from n down to 2, apply `RemoveInconsistent(Parent(X_j), X_j)`
 3. For j from 1 to n , assign X_j consistently with `Parent(X_j)`

Nearly tree-structured problems

- *Conditioning*: assign a value to a variable, prune neighbor's domains
 - Goal is to get a tree
- *Cutset conditioning*: for a subset, find all assignments that result in trees
 - Cutset cannot be too big
 - How might you choose the cutset?