

Informed Search

- Reading: Russell and Norvig, ch. 4.1-4.2
 - The material here is an edited version of Russell's slides with some different examples
- Best-first search
- A* search

Pseudocode for tree search

- Search strategy affects which node is expanded next

// edited to look a bit more object-oriented

function tree_search(problem, fringe) **returns** a solution or failure

fringe.insert(make_node(problem.initial_state()))

loop do

if fringe is empty **return** failure

 node ← remove_front(fringe)

if problem.goal_test(node.state()) **return** node

 fringe.insert_all(node.expand(problem))

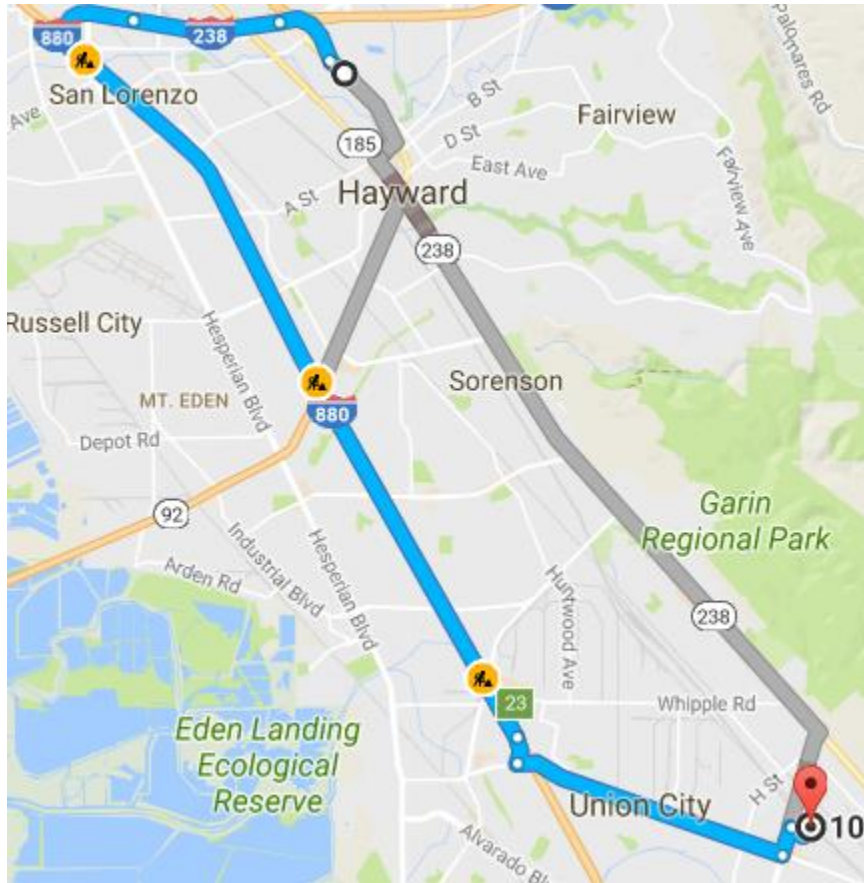
Best-first search

- Estimate how likely a node is to lead to the optimal solution using an evaluation function $h(n)$
 - “Estimate” – evaluation function is a *heuristic*
 - Implement fringe as priority queue ordered by evaluation function
 - Example of a greedy search
- Ex: for travel, distance between node’s location and destination (if you could fly between the 2 points)
 - Is the travel time computed by google maps exact?

Best-first search example

- Travel from 21500 Foothill Blvd, Hayward to Union City BART station (10 Union Square, Union City)
- Could first head towards DNA Motor Lab at 21739 Mission Blvd or 238 North on-ramp (close to 21300 Foothill Blvd)
 - If “best” choice is estimated from physical distance heuristic, choose DNA
 - This route only uses local roads, so it is actually slower (usually)
 - When would physical distance be a better heuristic?

Maps for best-first example



Best-first search

- Complete?
 - How might we end up in a loop? (Consider what might be the cause of this reported bug:
<https://steamcommunity.com/app/402310/discussions/2/405693392926963728/>)
 - Yes if in a finite state space and checking for duplicate states
- Time, space
 - Both $O(b^m)$ – may end up keeping all nodes in memory, choice of heuristic has a significant impact on time and space
- Not optimal

A* search

- Evaluation function $f(n) = g(n) + h(n)$
- $g(n)$ = cost to reach node n
 - reduce chance of expanding already expensive path
- $h(n)$ = estimated cost to reach node with goal state (“goal node”)
 - $h(n) \leq h^*(n)$, $h^*(n)$ = actual cost to reach goal node
 - This property makes $h(n)$ an *admissible* heuristic
 - $h(n) \geq 0$
 - $h(n) = 0$ for any goal node

Examples of admissible heuristics

- 8-puzzle
 - Number of tiles in the wrong position
 - Total “Manhattan” distance of tiles from their correct positions
 - Compute each for the following example from Gerhard Wickler (<http://www.aiai.ed.ac.uk/~gwickler/eightpuzzle-inf.html>), goal state on right

8		6
5	4	7
2	3	1

	1	2
3	4	5
6	7	8

Examples of admissible heuristics

- Number of tiles in the wrong position = 7
 - Only the 4 is in the right place
- Total Manhattan distance (for tiles 1, 2, 3, ...)
 - $f = 3 + 4 + 2 + 0 + 2 + 4 + 2 + 4 = 21$

8		6
5	4	7
2	3	1

	1	2
3	4	5
6	7	8

Dominance

- If heuristic $h_2(n) \geq h_1(n)$ for all n , then h_2 *dominates* h_1 and is better for search
- Time/space complexity is still exponential, but can be much faster/use much less memory
- Russell's sample numbers of nodes for the 8-puzzle problem:

Depth of solution	Iterative deepening	A* using # misplaced tiles	A* using Manhattan distance
14	3.5M	539	113
24	54G	39K	1.6K

Developing admissible heuristics


- Useful approach: find exact solution for relaxed version of problem
- For 8-puzzle:
 - If allowed to move tiles anywhere, # misplaced tiles is exact
 - If allowed to move tile to any adjacent square, Manhattan distance is exact
- Requirement: cost of optimal solution of relaxed problem \leq cost of optimal solution of real problem

Consistency

- A heuristic h is consistent if for any node n and its possible successor nodes n'
 - $h(n) \leq \text{cost of action to go to node } n' + h(n')$
 - So heuristic always underestimates cost
- If a heuristic is consistent, it will also be admissible, but not necessarily vice versa

Relaxed problem examples

- Another example: Traveling Salesman Problem – for N cities, find shortest sequence that visits each city exactly once
 - Classic NP-Complete problem – best known algorithm takes exponential time
 - Minimum spanning tree problem as relaxed version of problem
 - Selects edges of graph that connect all vertices with minimum cost
 - May not be possible to traverse edges without revisiting vertex



- Solvable in $O(n^2)$

A* example (uses Wickler's app for order of expansion)

- Start at node A: (screenshots use <http://mypuzzle.org/sliding>) ($f = 5$)

1		2
3	4	8
6	5	7

- Possible successor nodes B, C, D with states: ($f = 5, 7, 7$ – includes 1

	1	2
3	4	8
6	5	7

1	2	
3	4	8
6	5	7

1	4	2
3		8
6	5	7

A* example

- Expand B to get node E with state: ($f = 7$) (duplicates detected)

3	1	2
	4	8
6	5	7

- Expand this node to get nodes F, G: ($f = 9, 9$)

3	1	2
4		8
6	5	7

3	1	2
6	4	8
	5	7

A* example

- Expand C to get node H: ($f = 9$)

1	2	8
3	4	
6	5	7

- Expand D to get nodes I, J, K: ($f = 9, 7, 9$)

1	4	2
	3	8
6	5	7

1	4	2
3	5	8
6		7

1	4	2
3	8	
6	5	7

A* example

- Expand J to get nodes L. M: ($f = 9, 7$)

1	4	2	1	4	2
3	5	8	3	5	8
	6	7	6	7	

- Expand M to get node N: ($f = 7$)

1	4	2
3	5	
6	7	8

A* example

- Expand N to get nodes O, P: ($f = 7, 9$)


1	4	2	1	4
3		5	3	5
6	7	8	6	7

- Expand O to get nodes Q, R, S: ($f = 9, 9, 7$)

1	4	2	1	4	2	1		2
	3	5	3	7	5	3	4	5
6	7	8	6		8	6	7	8

A* example

- Expand S to get nodes T, U: ($f = 7, 9$)



- Node T has the goal state, so done!

A* search optimality

- Standard proof:
 - Suppose a suboptimal node G has been generated and is placed on the queue
 - G will not be expanded ahead of any node n on the path to the optimal solution node G_{opt}
 - $f(G) = g(G)$ because G is a goal node
 - $g(G) > g(G_{\text{opt}})$ because G is not optimal
 - $g(G_{\text{opt}}) \geq f(n)$ because h is admissible
 - So $f(G) > f(n)$ and n will be expanded before G

A* search optimality

- Alternative explanation
 - similar to breadth-first, but expands in “contours” with the same f value
 - In above example:
 - $f = 5$: nodes A, B
 - $f = 7$: nodes C, D, E, J, M, N, O, S, T
 - Not expanded: $f = 9$: nodes F, G, H, I, K, L, P, Q, R, U

A* search evaluation

- Complete
 - Yes, if the # nodes n where $f(n) < f(G)$, G = optimal goal node
- Time, space complexity: still exponential
 - Affected by accuracy of $h(n)$ and # nodes with sufficiently low $f(n)$
- Optimal: yes, A* expands
 - All nodes where $f(n) < C^*$, $C^* = f(G)$
 - Some nodes where $f(n) = C^*$,
 - No nodes where $f(n) > C^*$,