

Problem Solving and Search

- Reading: Russell and Norvig, ch. 3
 - The material here is based on Russell's slides as well as a version of the slides at UC Irvine: <https://www.ics.uci.edu/~rickl/courses/cs-171/cs171-lecture-slides/cs-171-02-IntroSearch.pdf>
- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms

Problem-solving agents (code as function)

```
function Simple-Problem-Solving-Agent (percept) returns  
an action
```

- **Decision-making = choosing an action**

```
  static:
```

```
    seq, an action sequence, initially empty
```

- **Generate plan to reach goal as sequence of actions**

```
  state, some description of the current world state
```

- **state represents agent + environment**

```
  goal, a goal, initially null
```

- **What agent should try to accomplish**

```
  problem, a problem formulation
```

Problem-solving agents (code as function)

```
state ← Update-State(state, percept)
```

- **Next state based on current state and input**

```
if seq is empty then
```

```
    goal ← Formulate-Goal(state)
```

- **Goal determined by environment**

```
    problem ← Formulate-Problem(state, goal)
```

```
    seq ← Search( problem)
```

- **Search for not just goal, but path to goal**

```
    action ← Recommendation(seq, state)
```

```
    seq ← Remainder(seq, state)
```

- **Taking an action affects plan**

```
    return action
```

Offline vs. online problem solving

- Above code is “offline”
 - Entire environment is known
- “online” problem solving = will be learning about environment during solution of problem

Example for offline solution

- Determine how to travel from one city to another
 - What makes an offline solution viable?
 - What real-world events might affect solution?
- Specifics: in Arad, Romania, flight leaves Bucharest tomorrow
 - State = Romanian road system, current city (in Arad)
 - Should probably consider other options like bus/train/taxi, but there may be other unmentioned factors, like needing to return a rental car
 - Problem makes location discrete: only make choices at cities
 - Formulate-Goal: get to Bucharest (in time to catch flight)
 - Formulate-Problem: agent's state = current city

Example for offline solution

- Formulate-Problem
 - Agent's state = current city
 - Action = drive from one city to another
- Solution/Plan = sequence of cities
- Graph from text is online in many places – ex:
<http://robotics.cs.tamu.edu/dshell/cs420/images/map2.gif>

Problem types

- Can model this problem several ways:
 - Deterministic, fully observable (“single state”)
 - Agent is always in known city, generate solution as sequence of cities
 - As UCI slides point out, if this is the case, Dijkstra’s algorithm can be used
 - Can also see situation as not capable of handling random events
 - Non-observable
 - Do know the possible states, but not which is the current state
 - Agent may lack sensors, knowledge of its own location, must still find solution
 - Ex: agent loses GPS signal and is moved
 - May be more useful to think of Roomba (automated vacuum cleaning)

Problem types

- Can model this problem several ways:
 - Non-deterministic, partially observable
 - Plan based on current knowledge
 - Update plan when new percepts update knowledge
 - Typical sequence: plan, one action, update plan, one action, update plan, ...
 - Unknown state space
 - Differences from non-observable
 - Can discover environment (sensors)
 - “online”

Vacuuming 2 squares

- Single state
 - Based on which squares are dirty (agent knows), vacuum current square if needed, then move to other square and vacuum
- Non-observable
 - Vacuum both squares
 - Goal is to have both squares clean, but you cannot tell which squares are clean, so vacuum just in case

Vacuuming 2 squares

- Nondeterministic
 - Can tell if necessary to vacuum current square, but may need to move to other square to check (dust does accumulate after a while)
- Unknown state space
 - Not applicable here
 - Compare to starting up vacuum robot in new house


Problem formulation

- Single state problems – need to define:
 - Initial state
 - Successor function
 - Set of (action, resulting state)
 - Again, possible because of extent of knowledge
 - Goal test
 - How does agent know if it's done?
 - Path cost
 - Cost of sequence of actions
 - Should be additive, generally nonnegative
- Solution = sequence of actions to get from initial state to a goal state

Modeling the real world

- Too many states, attributes to feasibly represent
 - Must abstract – represent many real world states with each abstract state
 - Vacuum example treats area as 2 positions
 - ☺ represents happy face for any person
 - Abstract action = combination of real actions
 - For vacuum to move, series of electrical and mechanical actions
 - For travel from one city to another, must travel along several roads
 - To guarantee plan would work in real world, action must work for any real state → must end in some real state modeled by abstract resulting state
 - Should be simpler than original problem (supports divide and conquer)
 - Abstract solution corresponds to set of sequences of real actions

Vacuum world state space (scan of textbook figure)





Vacuum world state space

- State = (robot location, [dirt in left pos?, dirt in right pos?])
 - Robot can only be in left or right box
 - Amount of dirt is not measured
- Actions = {L = left, R = right, S = suck, NoOp}
 - NoOp's are not in this diagram – how might they fit in?
- Goal test = no dirt, so either bottom state is a goal state
- Path cost = 1 per actions L, R, S
 - = 0 for NoOp

8-puzzle

- From <http://www.aiai.ed.ac.uk/~gwickler/images/8-puzzle-states.png>
- Left is a possible start state, right is a possible goal state



8-puzzle


- <http://www.cs.cmu.edu/afs/andrew/course/15/381-f08/www/lectures/HandoutUninformedSearch.pdf> notes following:
 - On average, 22 moves enough to reach solution
 - But 1.8×10^5 states (15-puzzle has 1.3×10^{12})
 - Representing every state is probably not desirable (though feasible)
- Note that $1.8 \times 10^5 < 9!$: https://en.wikipedia.org/wiki/15_puzzle notes that Johnson & Story (1879 – this is a very old problem) proved half of the possible states cannot reach the goal state

8-puzzle

- States: the values (1-8, blank) in some fixed order
 - Ex: 8, 0, 6, 5, 4, 7, 2, 3, 1 (with 0 for blank)
 - Only model position after each move is complete
- Actions: based on “moving” blank left, right, up, down
- Goal test: if goal state has been reached
- Path cost = 1 per move
 - Finding optimal path is NP-hard

Robotic assembly

- From text: <http://chalmersgu-ai-course.github.io/AI-lecture-slides/img/stanford-arm+blocks.png>



Robotic assembly

- States: angles of robotic arm joints (need floating point precision) plus parts to be assembled (need positions, maybe orientation)
- Initial state: initial angles, positions/orientations
- Actions: movement of joints
- Goal test: if parts are correctly assembled (and may want joints at specific angles to prepare for next assembly/avoid contact)
- Path cost: time to assemble (UCI slides point out energy cost may need to be included)

Tree search

- Each node stores a state
 - Root based on initial state
- Children nodes are states reachable by an action
 - Search tree by *expanding* leaf node to generate new nodes (states)
 - How is this tree different from a binary search tree? -- Number of children?
How much of tree is known?
 - Can often end up returning to same state
- Each node can be evaluated to see if it is a goal state
 - Search fails if no more nodes to expand and no goal state found
 - Can continue after a goal state is found to try to find better solution

Duplicate states

- Not always feasible to store all visited states to avoid expanding duplicate states
 - Ex: (from UCI slides), even 8-puzzle has $9!$ states = 362,880
- Feasible option, store states along path to root (in hash table)

Search Strategies

- Differ in choices of order of expanding nodes
- Can be compared on:
 - *Completeness* – does it always find a solution node if one exists?
 - Time complexity – biggest factor is number of nodes generated
 - Space complexity – number of nodes can be infinite, fitting nodes in memory also a concern
 - Optimality – guaranteed to find best solution?
- How to measure time and space complexity
 - Max branching factor = number of children
 - Depth of least cost solution
 - Max depth of state space